

Binary- and Multi- Class Logistic Regression

Outline

- ▶ Logistic Regression with Sparse Regularization
- ▶ Multi-Class Logistic

Logistic Regression

1. Load the binary dataset

- ▶ $X \in \mathbb{R}^{n \times d}$, $y \in \{\pm 1\}^n$

2. Standardize the columns of X

- ▶ Each column in X is standardized using z-score scaling
- ▶ $X^j = \frac{(X^j - \mu_j)}{\sigma_j}$, $\mu_j = np.sqrt{\frac{\sum_{i=1}^n X_i^j}{n}}$, $\sigma_j = \frac{\sum_{i=1}^n (X_i^j - \mu_j)^2}{n}$
- ▶ Each column has mean 0 and variance 1 but different ranges

3. Add bias (add a column of all ones to X)

4. Train logistic regression on X and y

Trained Logistic Regression

1. Standardize the columns of X_{validate} using μ_j and σ_j
2. Add bias (add a column of all ones to X)
3. Compute validation error on X_{valid} and y_{valid}
4. Compute the number of non-zeros for w

Fit logistic Regression

- ▶ Initialize $w \in \{0\}^d$
- ▶ Check gradient using the complex step method (it's given)
 - ▶ more accurate than finite differencing
- ▶ Set up the objective function and the gradient funObj (func, grad)
 - ▶ $\text{funObj}[0] = f(w) = \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i))$
 - ▶ $\text{funObj}[1] = \nabla f(w) = ?$ (Derive)
- ▶ Run findMin
 - ▶ Update w^t for **maxEvals** iterations
 - ▶ $w^{t+1} = w^t - \alpha \nabla f(w^t)$
 - ▶ Uses line search to determine α that ensures
 - ▶ $f(w^{t+1}) < f(w^t)$
 - ▶ $w^{t+1} = w^t - \alpha \nabla f(w^t)$

Evaluate logistic Regression

- ▶ Compute the classification error on the validation set
- ▶ Compute the number of non-zeros on w

Logistic Regression L2

- ▶ Logistic regression

$$f(w) = \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i))$$

- ▶ L2 Regularized Logistic Regression
- ▶ `funObj[0] =`

$$f(w) = \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i)) + \frac{\lambda}{2} \|w\|_2^2$$

- ▶ `funObj[1] = $\nabla f(w)$ = ?`

Logistic Regression L1

- ▶ $f(w) = \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i)) + \lambda \|w\|_1$
- ▶ $\|w\|_1$ is not differentiable at the origin

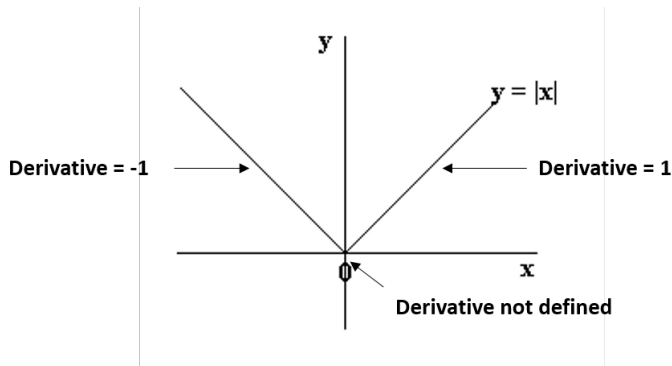


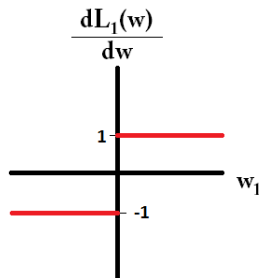
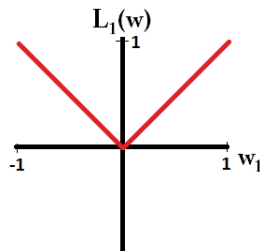
Figure 1:

Logistic Regression L1

- ▶ $f(w) = \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i)) + \lambda \|w\|_1$
- ▶ $\|w\|_1$ is not differentiable at the origin
- ▶ Define funObj the same way as Logistic regression
 - ▶ $\text{funObj}[0] = f(w) = \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i))$
 - ▶ $\text{funObj}[1] = \nabla f(w) = ?$
- ▶ Use findMinL1 instead of findMin to optimize the objective function
- ▶ findMinL1 uses proximal gradient to take care of the L1 norm

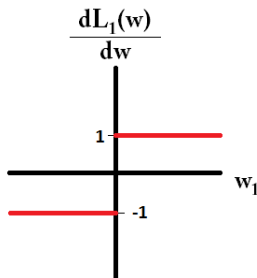
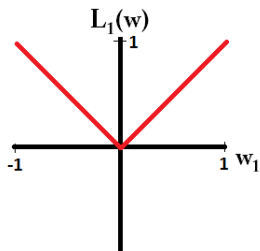
Logistic Regression L1

- ▶ You should expect some zero elements in w the parameter vector
- ▶ Distance to minimum is not negligible


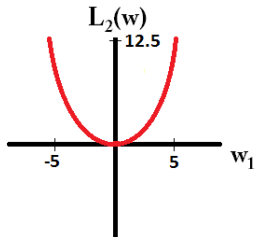


Logistic Regression L1

- ▶ Minimize $L_1(w) = \sum_i |w_i|_1$

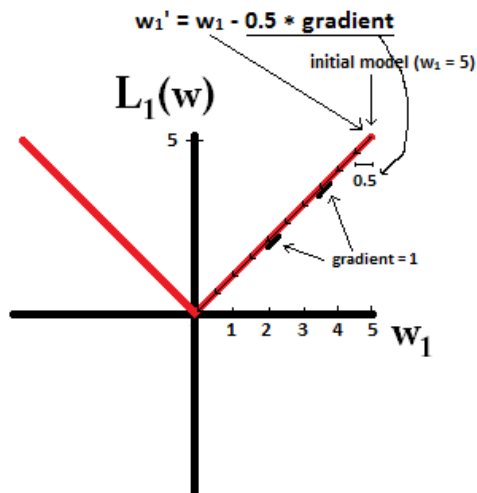


- ▶ Minimize $L_2(w) = \frac{1}{2} \sum_i w_i^2$



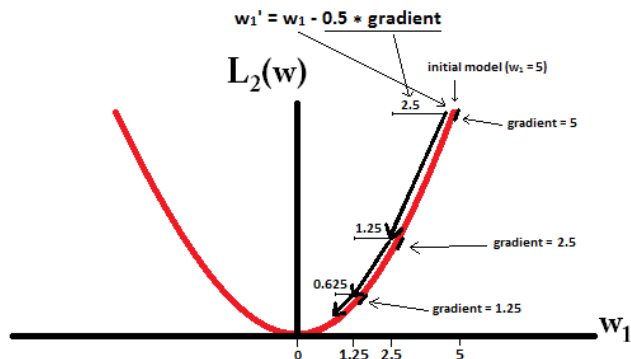
Logistic Regression L1

- ▶ Minimize $L_1(w) = \sum_i |w_i|_1$



Logistic Regression L1

- Minimize $L_2(w) = \frac{1}{2} \sum_i w_i^2$



Logistic Regression L0

- ▶ Objective:

$$f(w) = \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i)) + \lambda \|w\|_0$$

- ▶ Class inheritance (Demonstrate)
- ▶ Always include the bias
- ▶ Select best feature using forward selection
- ▶ Note that adding any feature adds the same penalty λ
- ▶ Stop when the decrease in the model error is less than the added penalty

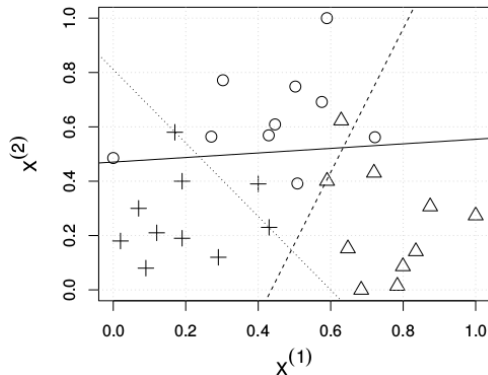
leastSquaresClassifier

- ▶ Classify a multi-class dataset using one-vs-all method
- ▶ Consider a dataset of 3 samples
- ▶ How many linear models ?

leastSquaresClassifier

- ▶ Classify a multi-class dataset using one-vs-all method
- ▶ Consider a dataset of 3 samples

Multiple binary classifiers



leastSquaresClassifier

- ▶ To make a prediction:
 - ▶ Compute $X_i W^k$ for each class k and take the largest.

```
def predict(self, X):  
    yhat = np.dot(X, self.W)  
  
    return np.argmax(yhat, axis=1)
```

logisticRegression

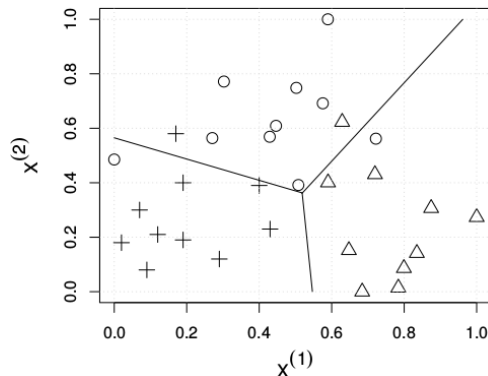
- ▶ Instead of leastSquares use logisticRegression
- ▶ $f(w) = \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i))$

```
for i in range(self.n_classes):  
    ytmp = y.copy().astype(float)  
    ytmp[y==i] = 1  
    ytmp[y!=i] = -1  
  
    self.W[:, i] = np.linalg.lstsq(np.dot(X.T, X),  
                                    np.dot(X.T, ytmp))
```

Multinomial logistic regression - SoftmaxClassifier

- ▶ Uses one classifier to classify points w.r.t all classes
- ▶ Learns a global distribution
 - ▶ Probability that a data point belongs to each class

Multinomial classifier



Multinomial logistic regression - SoftmaxClassifier

- ▶ Same as with the other objective functions
 1. define funObj (func, grad)
 2. use findMin to minimize the function
- ▶ Set up the objective function and the gradient funObj (func, grad)
 - ▶ funObj[0] =

$$f(W) = \sum_{i=1}^n \left[-w_{y_i}^T x_i + \log \left(\sum_{c'=1}^k \exp(w_{c'}^T x_i) \right) \right].$$

- ▶ funObj[1] = $\nabla f(w) = ?$
- ▶ Compute validation error